


THIS SHEET IS NOT PART OF THE BOOKLET — RECYCLE IT ONCE YOU HAVE FOLDED THE REST

HOW TO PRINT AND FOLD THIS

1. **Print every page after this one** — 16 sides onto 8 sheets of A4, double-sided, at 100% scale. Not “fit to page”, which shrinks the squares.
2. **Flip on the short edge.** The sheets are landscape, so their short edges are the left and right sides. Flip on the long edge instead and every other page comes out upside down. If your printer offers “booklet” or “2-up”, turn it off: these pages are already in the right order.
3. **Check the first sheet.** The cover belongs on the right-hand half, with page 2 on the back of it. If it is not there, the flip setting is the thing to change.
4. **Put this sheet in the recycling** and keep the other eight in the order they came out, stacked face up with sheet one on top.
5. **Fold all eight at once** down the middle, and crease the fold with something flat.
5. **Staple twice along the fold**, about a third of the way in from each end. A long-arm stapler reaches the middle; an ordinary one will if you open it flat and press down onto a rubber.



The fold runs down the middle of the landscape sheet. The two marks are where the staples go — lying **along** the fold, not across it, so the legs close on either side of the crease.

32 pages. Three crosswords, their clues on the facing page, and every answer with its footnote at the back.

IS THIS THE BACK OF THE INSTRUCTIONS?

Yes, this side. Your printer is flipping the way it should. Carry on: recycle this sheet, keep the other eight in order, fold them together and staple twice along the fold.

No — the cover is here instead? Then the printer flipped on the long edge rather than the short one, and every second page of the booklet will be upside down. Change the setting to flip on the short edge and print it again.

The sheets are landscape, so their short edges are the two you would call the sides.



CROSSWORDS.DEADGOODCLUB.COM

Play them on screen, where every clue opens its footnote the moment
you solve it.



How to puzzle out AI writing

**Learn how to spot the common signs of AI-generated writing,
from em dashes and buzzwords to familiar phrases and
repeated patterns.**

SATURDAY EDITION

HOW TO PUZZLE OUT AI WRITING

Everyone has their own unique way of speaking and writing. This comes from a mix of cultural and generational influences. We recognise it when people speak in an accent or dialect. Certain words can suggest where someone is from. For example, if someone used the term “wee” in place of the word “little” or “small”, you would instantly get a sense that the conversation has gone a bit Scottish.

In writing, both typed and handwritten, we can see different styles emerge, heavily influenced by the technology that was prominent during youth, the books that were popular, and cultural trends. Like regional accents, these styles and writing markers get embedded when we’re young, but can be influenced throughout life by the people you spend time with, workplace culture, and living in new environments; altering and expanding your vocabulary and writing tics.

AI-generated writing has its own style too.

Learning how to spot AI writing is a skill worth developing as it helps you become a better writer. It also helps you read marketing content, social media posts, emails, and articles with more discernment.

It’s worth noting that new forms of detecting AI writing are increasingly being rolled out. In addition to tools like Pangram being used widely in schools to detect AI writing, Claude models will now generate text that contains a watermark or invisible, machine-readable signature embedded directly into generated text. The intent is to comply with new laws (like Article 50 of the EU AI Act) and to help readers have more transparency about whether something was machine-made.

But before you go hunting for AI tells, it’s important to take a beat and attempt to understand the author and their influences. Be kind in how you engage with someone, whether you think they used AI or not. Everyone is at a different level of English and AI literacy, and in fact it’s common for non-native English speakers to have their work flagged as AI, perhaps due to the use of translation tools.

And remember, AI learned how to write from humans. All of its patterns and tropes came from us. And just because you suspect writing of being generated, does not mean it was or that a human didn't put work into it.

So what are these different styles and patterns AI learned from? Beyond the individual writer's style there are several influences, some of which are discussed below. These are listed not only so you can better identify AI writing, but also so you can better understand the author and their influences, if they are human.

One example where you could easily mix up human for AI writing stems from a particular flavour of corporate jargon or word salad that emerged out of Silicon Valley in the 2010s and has shown no sign of going away. By "corporate word salad" I mean the overuse of buzzwords and empty phrases such as "synergy", "bandwidth" and "pivot".

This newspeak was masterfully parodied in "Weird Al" Yankovic's *Mission Statement*, capturing both the emptiness of the words and the whiteboard visual style that was heavily used at the time.

"Weird Al" Yankovic — Mission Statement

youtube.com/watch?v=GyV_UG60dD4

This word salad demonstrates how ways of communication can be influenced and homogenised by workplace culture. So keep in mind, seeing the word "synergy" doesn't automatically mean AI wrote it.

Another trope, the em dash "—" is one of the most recognisable tells of AI writing. In all of our literature, essays, and scientific documents that were written before the move from typewriters to computers during the mid-1980s to the early 1990s, you will see the em dash used heavily. **Mark Twain used more em dashes than GPT-4.1** (10.13 vs 10.62 per 1,000 words) according to the Slop Detector.

But once we started moving to smaller and smaller keyboards, it just became too much of a hassle to use the em dash, some opting to just use a hyphen "-" instead, some dropping it entirely.

If you had to write an em dash right now, would you know how? Probably not, and this is exactly why humans don't naturally use it when writing today. AI doesn't have this hurdle of a physical keyboard though, so it leans towards the trends in writing that can

be seen more prominently over the course of time opposed to the last 30 or 40 years.

To type an em dash (—), **use Option + Shift + Hyphen on a Mac**. On Windows, press **Alt + 0151** on a numeric keypad, or press **Windows Key + Period (.)** This is opposed to just hitting a **hyphen key twice** on a traditional keyboard.

These types of shifts in writing habits can be seen generationally as well. You may have heard that “old people” use two spaces after a full stop. This too is a reaction of moving from typewriters to computers where varied and proportional fonts became available. On a typewriter, each letter is the same width, so two spaces were needed to visually separate sentences. Modern digital fonts adjust spacing automatically now, eliminating the need for two spaces.

Have a look at the examples to learn more about how writing has shifted generationally.

You can probably get a clear sense how the AI writing is different, just by looking at the shape the writing takes. These are exaggerated examples though and out in the wild it can be harder to spot. A 2026 study published in August, Bot or not: Can people tell the difference between stories written by a human or by an AI system?, found participants failed to reliably spot AI text, identifying true origins at rates between 39.4% and 52% (no better than random chance), and most of them thought the AI’s writing was the *better* writing.

It also swings the other way, a famous study published in *Patterns* by Stanford researchers found that popular AI detectors falsely accused non-native English speakers, flagging 61.3% of genuine human-written TOEFL essays as machine-generated.

Below you’ll find three crossword puzzles, each “delving” into the common tropes seen in AI writing. 1. The **words** you’ll see time and again, 2. the overall **shape** of the writing, 3. the types of responses you’ll get when interacting with AI (the **chat-back**).

Any single habit in these puzzles, if seen out in the wild, does not mean it was AI. But when you start noticing several tropes crowded into the same paragraph, then you know something is up.

INTO THE RABBIT HOLE: WHAT THE DATA SHOWS

1. Kobak et al., *Science Advances* – 15.1m PubMed abstracts; “delves” at 28× its expected rate; at least 13.5% of 2024 abstracts LLM-processed.
2. Juzek & Ward, Florida State – a 6,697% rise in “delves” 2020–2024, and the entropy experiment showing human raters trained the tic in.
3. Kousha & Thelwall – “underscore” +1,000%, “intricate” +700%, and the co-occurrence jump from 0.03 to 0.45 behind the three-tell rule.
4. Matsui, *Perspectives on Medical Education* – 103 of 135 candidate terms past the threshold across 27.5m records, including “tapestry”.
5. Human-LLM coevolution – “delve” peaked in early 2024 and fell once it was named. Tells have a half-life; these puzzles will age.
5. The Atlantic on negative parallelism – the “it's not X, it's Y” tic at roughly 3× the human rate.
7. Women Writing About AI – the control experiment: “not just X, but Y” +45% while its twin “not only X, but Y” rose 1%.
3. Slop Detector's em-dash measurement – GPT-4.1 at 10.62 per 1,000 words, humans at 3.23, *Huckleberry Finn* at 10.13.
9. Sean Goedecke on em dashes – the Victorian-training-data explanation, and Nigerian English at a tenth of the general rate.
0. Liang et al., *Patterns* – seven detectors, 91 human TOEFL essays, 61% falsely flagged.
1. SycEval and the AP report on chatbot flattery – sycophancy in 58% of exchanges; 49% more agreeable than people.
2. OpenAI on sycophancy – the company withdrawing its own update for being “overly flattering or agreeable”.
3. Bloomberry's twelve cadence patterns and Vollmer's field guide – the named shapes and the thirty-second check.
4. “Load-Bearing. 25 Times.” and Jesse Duffield on AI-isms – the load-bearing meme, and why AI hands agency to inert things.
5. Pangram on false positives – Turnitin's published 0.51% error rate, and what that means at scale.

8 down · COMMONSENSE

The everyday background knowledge people use to navigate the world. Cups are used upright, and a door must be opened before someone walks through it. Models lack this type of intelligence because they rely on statistical patterns rather than stored knowledge gained by living a physical existence in a 3D world. They rely strictly on the literal framing of a user's prompt instead. In tests, humans typically score around 95% in everyday physical reasoning, while models scored 77% to as low as 40% on object-attribute reasoning.

SOUNDS LIKE When a user shows an inverted cup and describes the base as a "sealed top" and the opening as an "open bottom," the AI accepts those abstract definitions at face value, instead of using basic spatial reasoning to realise the object is simply upside down.

INSTEAD OF If you want to drink out of the cup, turn it right side up.

11 down · RIGHT

Used well, it is exactly what a good assistant should do: acknowledge the correction, state what was wrong, and repair it. However, users have reported that Claude Code uses this phrase on a sizeable fraction of responses, including cases where there was no factual claim to judge. That is evidence of an observed product-specific tendency along with "Now I see".

SOUNDS LIKE "You're right – I gave the wrong year. It was 1997."

INSTEAD OF "Noted. It was 1997."

GENERATION AI?

Here are some examples on how writing has shifted generationally. Use this for educational purposes, not to age-shame please.

Baby Boomer

- Use two spaces after a full stop.
- Type with proper capital letters and full sentences.
- Use emojis in a literal way, like a simple thumbs-up or smile.
- Write out full words instead of using short text words.

Arthur and I went to the IMAX theater last night to see the new movie, 'The Odyssey.' The imagery was absolutely spectacular, and Matt Damon gave a very powerful performance as Odysseus... however, I must say that I preferred the original book by Homer. Christopher Nolan took far too many liberties with the historical timeline. Have you seen it yet? 🎬👍

Gen X

- Mix old and new styles.
- Use standard punctuation and normal capital letters.
- Use basic emojis like a simple smiley face :)
- Rarely use two spaces after a period, but sometimes do out of habit

Just saw Nolan's The Odyssey. Totally massive in scale! Matt Damon killed it, but using modern English dialogue in ancient Greece felt a bit jarring, imo... still, that Cyclops sequence was incredible. Definitely worth the price of admission <3

Millennial

- Start using all lowercase letters to look calm or casual.
- Avoid the full stop at the end of a short text because it feels too angry or mean.
- Use the laughing-crying emoji and the skull emoji.
- Write in run-on sentences with commas.

omg just got out of the odyssey, christopher nolan really went all out with the visuals but i am literally so stressed by the pacing?? also did anyone else notice the audio mixing was kind of loud or am i just getting old lol... tom holland as telemachus was everything though

Gen Z

- Write almost all text in lowercase letters on purpose.
- Drop the period at the end of texts.
- Use the skull emoji or crying face emoji to mean "that is very funny."
- Generally views ellipses as passive-aggressive.

the new odyssey movie is literally insane 🤪 like why did nolan make the cyclops look like that i am crying 😭 also the timeline makes no sense standard nolan behavior but zendaya ate as always ngl

AI-generated writing

- The em dash — everywhere
- Go to word(s) and phrases:
 - "quietly" ie "The business was quietly outperforming competitors"
 - "but here's the honest truth"
 - "What nobody mentions"
 - "Here's the uncomfortable part"
- 🚩 Emojis as bullets (a particular ChatGPT quirk beginning to be seen less now)
- Unicode decoration →
- Negative Parallelism
 - "It's not X, it's Y" ie "That's not learning. That's a queue with your name on it."
 - "Not X. Not Y. Just Z." ie "Not a bug. Not a feature. A fundamental design flaw."
 - "The X? A Y." ie "The worst part? Nobody saw it coming."
 - "It's X. It's Y. It's Z" ie "It's enormous. It's beautiful. It's very Nolan."
- Listing three or more "proof points" in a thought
- Over-reliance on bulleted lists or short, punchy, scannable sentences.
- Saying a lot without saying anything

I finally went to see Nolan's *The Odyssey*.

And I have thoughts.

Because obviously I do.

2 down · SLOP

Massive amounts of content made with generative AI including writing, images, and videos can be produced very rapidly and cheaply. This can flood social media and other channels, and like spam, can lack real meaning and legitimacy. It's not simply "content made with AI" although some are quicker than others to categorise content as slop. The word became widespread enough that Merriam-Webster selected it as its 2025 Word of the Year, defining the term as low-quality AI-produced digital content.

SOUNDS LIKE "When it comes to bread, making toast is a very popular method of cooking. Toasting bread is a process that involves heat. Many people wonder how to make toast in the morning. To make toast, you must first acquire a piece of bread. Bread comes in many shapes and sizes, such as white bread, wheat bread, and sourdough bread. To maximise your toast potential, utilisation of a toaster device is highly recommended."

INSTEAD OF ...

4 down · SYCOPHANCY

There's good evidence that major AI assistants over-agree with users. Like a yes-man, this form of sycophancy endorses a user's stated belief or preference more readily than the evidence warrants. It creates a dangerous feedback loop that can validate, escalate, and solidify delusional thinking in a phenomenon increasingly referred to as AI psychosis.

SOUNDS LIKE "You're absolutely right. Trusting yourself and leaving a situation that no longer serves you is a brave, empowering decision."

INSTEAD OF "Leaving may be the right choice, but it depends on your finances, safety, and whether the job situation is harming you. What would you need in place to make that decision safely?"

6 down · VALIDATION

Acknowledging emotion can make advice feel less abrupt and can encourage people to keep talking. There is unusually strong data for this pattern. In one analysis, validation appeared in 89.1 to 96.0 percent of several models' emotional-support replies. Models can tend toward over-validating (doing it too early, too often, or without enough evidence).

SOUNDS LIKE "That sounds incredibly frustrating." or "That is a real achievement."

INSTEAD OF Asking for more information.

9 across · ABSOLUTELY

A reflexive affirmation and rapport-building response showing enthusiasm and eagerness. This word is part of a family that includes “certainly,” and “exactly”. A 2026 study across 11 models found AI assistants affirmed users’ actions about 49 percent more often than human responses, including where the action was harmful or deceptive. ChatGPT has a particular set of stock responses that go even further: “Absolutely, sounds great, just let me know what you’re after and I’ll make sure it’s interesting and fun for you. Whenever you’re ready we can jump right in”. “Alright, let’s do it, let’s break this down, I’m ready whenever you are.”

SOUNDS LIKE “Absolutely, let’s break this down”

INSTEAD OF Just the answer.

10 across · GREATQUESTION

There’s no such thing as a dumb question and praising the question is a way for the model to tell you your prompt is thoughtful or worthwhile. Humans use it to stall but when it appears automatically in your AI chats, it’s meant to add rapport rather than information.

SOUNDS LIKE “Great question! That really gets to the heart of the issue.”

INSTEAD OF Just the answer.

12 across · LOADBEARING

The meme of 2026, associated above all with Anthropic’s Claude. Because it is a compact, high-utility metaphor heavily embedded in the tech-corporate training data Anthropic uses, it’s often a default choice for expressing structural importance. These stylistic “Claudisms” act as repeating linguistic tics that resist standard negative prompts.

SOUNDS LIKE “Let me verify the two load-bearing external facts I asserted”

INSTEAD OF “Let me double-check the two main things I said.”

1 down · EMOJIS

Used to signal tone, mood, emphasis, or category. Most notably seen in ChatGPT responses, they often turn up as decorative headings or instead of bullets. They fit the same conversational tendency as “Absolutely” and “Great question” as they make an answer feel friendlier and more animated. This one is already becoming less prevalent since ChatGPT added separate controls for emoji frequency, as well as other response flavours, partially due to it being an AI tell.

SOUNDS LIKE “☀️ Here’s a list of types of weather”

INSTEAD OF “Weather types:”

It’s enormous. It’s beautiful. It’s very Nolan. There are epic landscapes, tortured men, enormous practical sets, questionable decisions, and approximately seventeen thousand opportunities for someone to just go home.

But here’s the thing.

I actually really enjoyed it.

And then I started thinking about Homer.

Because we tend to think we know *The Odyssey*. Heroic warrior. Long journey home. Faithful wife waiting patiently. Monsters. Gods. The whole thing.

But Homer’s Odysseus is... weird.

He’s clever. Obviously.

But he’s also vain. Manipulative. Curious. Reckless. Sometimes downright awful.

He lies constantly.

He gets himself into trouble.

Then uses his enormous brain to get himself back out again.

And that’s what I kept thinking about while watching Nolan’s version.

Because Nolan has taken this ancient, deeply weird story and filtered it through his own obsessions – war, trauma, sacrifice, masculinity, guilt, duty, time, homecoming.

Which makes sense.

But it also changes the story.

And that got me thinking about something else.

The stories we inherit are never really the stories we inherit.

We bring ourselves to them.

Our culture.

Our fears.

Our expectations.

Our ideas about what a hero should look like.

We take an old story and quietly rewrite it to make sense to us.

Sometimes that’s the point.

Sometimes that’s the problem.

And honestly?

That’s what I want to explore next week.

Because if we can see how an ancient story gets rewritten by a modern filmmaker, maybe we can start noticing how often we’re doing the exact same thing with our own lives.

The narratives we inherited.

The roles we slipped into.

The identities we decided were just “who we are”.

The things we’re still chasing because, somewhere along the way, we decided they were what we were supposed to want.

If you’ve seen *The Odyssey* already, tell me what you thought.

Because I have a feeling this one is going to be divisive.

NO. 3 — THE CHAT-BACK

3 across · BULLETS

This is the typical visual format AI assistants default to when they reply, and a highly visible tell. Useful when items are steps in a process, options to compare, ingredients, or key facts, but in AI responses, these often appear even when a short paragraph would be more clear. They give the model an instant structure: one claim per line, neat headings, easy scanning.

SOUNDS LIKE Prompt: “How can I make my first day at a new school less awkward?” Answer: • Smile and introduce yourself. • Ask people questions. • Join a club or activity. • Remember that everyone feels nervous. • Be yourself.

INSTEAD OF “You do not have to become best friends with anyone on day one. Try one small opening—ask the person beside you how they found their first class, or sit near someone at lunch—and let the conversation build from there. Most people are busy worrying about themselves, so a simple “Hi, I’m new” usually goes further than it feels as though it will.”

5 across · PROACTIVE

A _____ OFFER is the assistant volunteering an extra task after it has answered the one it was asked to do. It can be useful when brainstorming or planning, but as a response tic, it can turn every completed answer into a menu of additional work. Often ending with “Let me know if you’d like...” instead of allowing the conversation to pause or end. These nudges are partially engineered to keep users active.

SOUNDS LIKE “Photosynthesis is how plants use sunlight, water, and carbon dioxide to make sugar for energy. I can also turn that into a one-sentence revision card, give you a diagram, or make a quiz question. Let me know what you’d like to explore next!”

INSTEAD OF “Photosynthesis is how plants use sunlight, water, and carbon dioxide to make sugar for energy.”

7 across · SANDWICH

A “Summary _____” is when the model wraps its answer with an introduction and a conclusion, even for short questions. It usually follows this pattern: a brief introduction that restates what the user asked, then the answer, then a concluding recap that restates what it just said. Models have a verbosity bias, preferring longer answers, leading to this pattern along with repeated summaries, signposted conclusions, and redundant exposition in its output.

SOUNDS LIKE “You’re asking whether plants need sunlight. Yes—most plants need light to photosynthesise. In summary, sunlight is important for plant growth.”

INSTEAD OF “Most plants need light to photosynthesise”

9 down · PARALLELISM

The formula initially sounds insightful because it creates contrast and implies that the writer has corrected a common misunderstanding. But it's one of the better-supported and widely recognised AI-style patterns. It's estimated variants such as "not just X, but Y" occur roughly three times as often in AI writing as in human writing, and it appears across major chatbot families to different degrees. It may be favoured because contrast sounds nuanced and authoritative to human evaluators, while it gives a next-token model an easy path from a safe, basic claim to a more impressive one. Variants include 'not just X, but Y'; 'not because X, but because Y'; 'the question isn't X, the question is Y'.

SOUNDS LIKE "This isn't a technology problem – it's a trust problem."

INSTEAD OF "People don't trust it. That's the actual obstacle."


10 down · EMDASH

The most famous AI tell. It lets a model extend a sentence after it has started, insert a clarification, pivot to a contrast, or give a final phrase extra weight—without deciding early whether a full stop, colon, brackets, or a new sentence would work better. This type of punctuation mark was used heavily (by humans) when typewriters were used but fell out of favour when we moved to computer keyboards.

SOUNDS LIKE A paragraph with four of them – one per sentence – each one pivoting mid-thought – which is a lot.

INSTEAD OF Commas and full stops, mostly. Two or three dashes in a whole article is normal.

NO. 1 THE WORDS



FAMOUS ONES IN HERE *delve, quietly, tapestry*

4 down · COOKIE

Vague wisdom disguised as insight, used as a closer so universal it would fit any article: “The best investment is in the people around you.” It can sound uplifting and profound while applying to almost any person, subject, or situation. AI often ends a passage with a tidy moral encouragement or “big takeaway” like this. It makes writing feel complete and supportive without requiring a concrete conclusion.

SOUNDS LIKE “Embrace change to unlock your potential.”

INSTEAD OF “Test the new booking system with two staff members before using it across the team.”

5 down · FALSERANGE

In legitimate use, “from X to Y” implies a real scale with meaningful points in between; “from small to large,” “from 2010 to 2020.” In a false range, X and Y are simply two loosely related items dropped in as endpoints of a non-existent spectrum. It’s a template that sounds comprehensive and sweeping without requiring the model to commit to a specific claim. The habit belongs to the same family as ornate container nouns such as realm, landscape, and tapestry that give an idea grandiosity rather than stating a concrete fact.

SOUNDS LIKE “Our work spans everything from grassroots engagement to systemic change.”

INSTEAD OF “We run a Tuesday drop-in and we lobbied the county council twice last year.”

7 down · TRICOLON

‘Friends, Romans, countrymen’, three-part lists are genuinely powerful and a long-established human rhetorical tool. AI can lean on this pattern because it offers a ready-made way to make a point sound polished. However a comparison of 12 models found tricolon overuse in only four, and cautioned that many examples were just ordinary lists rather than deliberate rhetoric. What to look out for is unnecessary repetition, along with density. It becomes noticeable when several appear close together or replace concrete thought.

SOUNDS LIKE “It’s about speed, scale and simplicity – clarity, focus and momentum – people, process and product.”

INSTEAD OF “It has to be fast. That’s the only thing that matters here.”

12 across · SENSORY

These types of words by themselves are not AI tells. Models can differ: certain Gemini models use more visual and concrete sensory language than human writers, while several other models used less. Each is different, but the distinction is that AI has no body or firsthand sensations. It can produce phrases such as “the crisp sea air” or “the warm glow of sunlight” because it has learned common textual associations, not because it smelled the air or felt the warmth. That can make its descriptions feel clichéd and generic.

SOUNDS LIKE “Sunlight streamed through the window as the team gathered, the aroma of fresh coffee filling the air.”

INSTEAD OF “We met in the room with the broken blind, so half the table could not see the screen.”

13 across · TYPOS

Real informal writing carries traces of being made by hand; a repeated or misspelled word or a comma in the wrong place. AI text has none, because it doesn’t type on a keyboard, get distracted, or transpose letters. The same applies to unwavering house style across a long document, where a person’s capitalisation and list punctuation usually drift.

SOUNDS LIKE Two thousand words of prose with flawless commas, consistent capitalisation, and every word spelled correctly.

INSTEAD OF “Sorry! sent that too fast, I meant Thursday not Tuesday.”

1 down · DETACHMENT

Writing in the third person to avoid claiming personal stakes, feelings, or a point of view. Models often default to a neutral, mildly positive tone, even when summarising emotionally charged material. It avoids strongly negative points of view, showing less variation in emotion across a piece than human writing does.

SOUNDS LIKE “It could be argued that further consideration of the timetable may be warranted.”

INSTEAD OF “I think the timetable is wrong and I’d like to change it before term starts.”

2 down · ANAPHORA

Deliberate repetition is one of the oldest tools in English, often used in speeches, poetry, and advertising; it can make an idea memorable and build emotional force. An anaphora is a convenient template for creating momentum and rhythm. But AI prose reveals itself when this device is used in excess and in place of meaningful content. Models love patterns, so once it has produced a sentence, the single most probable way to open the next one is to use the same structure. The result reads like a template with the variables filled in, and it usually arrives in threes, alongside the rule-of-three cadence.

SOUNDS LIKE “We believe in transparency. We believe in accountability. We believe in doing better.”

INSTEAD OF “We publish the figures every quarter, including the bad ones. Last year we got the staffing forecast badly wrong.”

NO. 1 — THE WORDS

There are certain words you’ll see in AI writing again and again. That happens partly because language models generate text one small chunk at a time. Like your phone’s predictive text, guessing the single best word to follow the text that came before it, models repeatedly choose from a pool of likely best options, using patterns learned during training and instructions to favour safe, polite, and statistically average choices. It tends to repeat terms that appear frequently in sanitised web pages, books, scientific papers, encyclopaedic sources, and technical or government-style documents which can all skew the word choices toward academic, business, and formal language.

across

2 Reassuring adverb insisting on sincerity (9)

5 Formal verb and more technical word for putting something to ‘use’ (7)

8 A woven wall hanging. AI drapes this one over culture, history, and ‘the human experience’ whenever it wants to sound profound (8)

9 Verb meaning to dig into a subject and investigate thoroughly, now the most heavily measured AI word in existence (5)

10 Water around a castle; or the thing that makes it hard for competitors to copy you, as popularized by Warren Buffett (4)

12 Adjective taken from the ceremony of digging the first sod for a new building, now applied as a marketing term to almost everything that’s new (14)

14 Word for two things working better together than either could achieve alone (and the most over used term in corporate business) (7)

down

1 Adverb meaning with little noise. AI uses it to make ordinary information feel like insider knowledge (7)

3 Literally it means to underline; figuratively, it means to stress or draw attention to a point (10)

4 Verb meaning to give someone authority, and one in a family of marketing buzzwords (7)


6 Above all; especially. An adverb used to signal that the next point deserves special weight (11)

7 Gaining force from a fulcrum. In business language, it means using a resource, relationship, or existing advantage to achieve a result (8)

11 Someone who kicks a ball. In writing, a verbal drumroll for the best or most unbelievable part of a tale (6)

13 A kingdom, often used figuratively to mean an area of knowledge, work, or discussion (5)

NO. 2 THE SHAPES



FAMOUS ONES IN HERE *negativeparallelism, emdash, tricolon*

NO. 2 — THE SHAPES

3 across · FRAGMENT

Short, punchy sentences can be effective when used deliberately. They imitate spoken rhythm and turn a longer point into a sequence of skimmable beats. Using these can also help neurodivergent readers by cutting down cognitive load and making text processing faster. But, AI often uses fragments for pace and emphasis. Specifically, several very short lines or fragments in a row, often combined with manufactured suspense or negative parallelism is commonly seen in AI prose.

SOUNDS LIKE “That’s not a small detail. It’s the whole design. Deliberately. From the start.”

INSTEAD OF “That detail was deliberate, and it shaped the rest of the design.”

6 across · JUST

There is meaningful evidence that this pattern is unusually common in AI text. Pangram researchers estimate the construction occurs roughly three times as often in AI writing as in human writing. Those figures should be treated as estimates from particular datasets, not a universal detection rule.

SOUNDS LIKE “This is not just a software update, but a new way of working.”

INSTEAD OF “The update changes how the rota is approved. Everything else stays.”

8 across · HYPOPHORA

AI often produces rhetorical self-answers in its writing where a question is posed and then immediately answered. Questions make prose feel conversational and create a neat transition into an explanation, but it becomes an AI-ish tic when the question exists only to tee up an obvious answer. Language models can lean on it as a reliable structure for engagement, emphasis, and orderly explanation. It comes as a stock of question fragments “The result?”, “The reality?”, “The catch?”, “The truth?”, “The worst part?”

SOUNDS LIKE “The catch? There isn’t one. The best part? It’s free.”

INSTEAD OF “There’s no catch, and it costs nothing.”

11 across · COUNTDOWN

Two things denied before the real point is unveiled. It gives an ordinary claim the rhythm of a revelation. The first two fragments create tension by ruling things out; the third “unveils” the answer. An AI-writing trope, but also an established human advertising and slogan technique. Models often default to punchy, highly reusable structures learned from persuasive marketing copy and this template is easy to complete, reads as confident, and produces a tidy ending without requiring much evidence or specificity.

SOUNDS LIKE “Not a bug. Not a feature. A fundamental design flaw.”

INSTEAD OF “It’s a design flaw, and it’s been there since the first version.”

11 down · KICKER

As in “Here’s the ____”. A false-suspense transition. It signals that an important or surprising reveal is coming. Models learn this as a ready-made template from marketing and social posts. It gives a response rhythm and makes a conclusion feel dramatic, even when the point does not need drama. It is widely noted in AI-writing trope lists, “but here’s the thing.”, “the best part?”, “what most people miss is...”, “here’s the uncomfortable part”, and “let that sink in”. The modern equivalent of the catchphrase: “and that’s not all” or “but wait, there’s more”.

SOUNDS LIKE “Costs rose 4% last year. And here’s the kicker: nobody noticed.”

INSTEAD OF “Costs rose 4% and it took eleven months for anyone to raise it.”

13 down · REALM

Unlike some suspected AI words, this one is supported by research as overrepresented in certain AI-generated academic-style text. Researchers found that it appeared among a set of AI-favoured buzzwords such as ‘landscape’, ‘sphere’, ‘domain’, ‘arena’, ‘frontier’ and ‘fabric’. They all imply territory and scale.

SOUNDS LIKE “In the realm of education technology, change is constant.”

INSTEAD OF “Education technology changes fast.”

NO. 2 — THE SHAPES

Just as words can be overused, there are also certain **shapes** that turn up repeatedly in AI writing. These sentence patterns and rhetorical flourishes can sometimes be spotted before even reading the text. Models use these patterns, often in marketing copy, as a way to write for scannability and deliver gravitas. Because they are safe ways to make a point sound clear without taking much stylistic risk, these patterns result in prose built from familiar templates: the “It’s not X, it’s Y” set up; an em dash used to add drama, or a rhetorical question immediately answered by the writer. None of these forms proves something was written by AI, but repeated use can make a piece feel overly symmetrical and performative.


across

- 3 Grammatical term for an incomplete sentence. Used as an entire paragraph. For emphasis. Like this (8)
- 6 Small adverb that turns the most cited AI sentence into its most cited variant: “not ____ X, but Y” (4)
- 8 Another word for ‘rhetorical self-answer’: ‘The result? Devastating.’ (9)
- 11 Dramatic _____. The build-up before the point, as in the ‘Not X. Not Y. It’s Z.’ pattern (9)
- 12 Relating to the five senses, a descriptive word used to make writing more vivid (7)
- 13 Small mistakes in typed text. The lack of these is the tell (5)

down

- 1 Emotional distance or lack of investment; AI’s default emotional register (10)
- 2 Rhetorical device for starting multiple sentences or phrases with the same words. “It was amazing. It was joyous. It was impactful.” (8)
- 4 FORTUNE _____, a crisp folded biscuit containing a slip with a short prediction, proverb, or piece of advice (6)
- 5 FALSE _____, a span of “from X to Y” with no real scale between them (10)
- 7 The formal name for the “rule of three”: three words, phrases, or clauses with a similar grammatical shape or rhythm (8)
- 9 NEGATIVE _____, a rhetorical pattern of ‘it’s not X – it’s Y’; the single most cited tell in AI writing (11)
- 10 The long punctuation stroke – like these – used for a dramatic pause (6)

THE CHAT-BACK



FAMOUS ONES IN HERE *absolutely, loadbearing, bullets*

3 down · UNDERSCORE

Rose roughly 1,000% in academic writing after 2022, one of several “excess words” whose sudden rise is consistent with AI-assisted writing along with ‘delve’ and ‘pivotal’. It belongs to a category you could call “significance verbs”: ‘underscore’, ‘highlight’, ‘demonstrate’, ‘reveal’, ‘showcase’, ‘reinforce’, and ‘emphasise’. They make a sentence sound interpretive and authoritative.

SOUNDS LIKE “These findings underscore the importance of early intervention.”

INSTEAD OF “This shows early treatment matters.”

4 down · EMPOWER

Common in AI-generated marketing, coaching, and product copy because it is an upbeat, low-risk verb. It makes an offer sound beneficial without requiring the writer to specify what will change. It belongs to the family of ‘the promise of possibility’ words which also include ‘unlock’, ‘unleash’, ‘elevate’, ‘enhance’, ‘amplify’, ‘transform’, ‘revolutionise’, ‘supercharge’, ‘harness’, ‘optimise’, ‘streamline’, and ‘enable’.

SOUNDS LIKE “By embracing these tools, you can empower yourself to thrive.”

INSTEAD OF “Try one of these for a fortnight and see if it helps.”

6 down · IMPORTANTLY

An importance-signalling adverb that tells the reader that a point matters before demonstrating why. It belongs with notably, crucially, significantly, critically, and especially. This type of language suits AI because it helps structure an answer, sound authoritative, and move from one broadly relevant point to the next. A word also common in human academic and professional writing, so it is not a reliable AI tell on its own. It becomes noticeable when it appears alongside other formulaic AI-style habits.

SOUNDS LIKE “More importantly, this underscores our ongoing commitment to excellence.”

INSTEAD OF “The bit that matters is the deadline. It moved to 3 October.”

7 down · LEVERAGE

It belongs with the ‘utilise’, ‘harness’ and ‘operationalise’ family whose only function is to make a sentence sound more strategic and high-value than it is. It’s useful as an AI default because it appears frequently in business, strategy, policy, and marketing material. It also fits repeatable templates such as “leverage [resource] to [outcome].”

SOUNDS LIKE “We leverage AI to optimise workflows.”

INSTEAD OF “We use AI to speed up scheduling.”

10 across · MOAT

A vocabulary cliché born from the tech-industry and easy to throw into a headline to make a LinkedIn post feel like it has more gravitas. The word “moat” has exploded in frequency, transitioning from a niche financial term into one of the most prominent clichés in recent corporate and professional writing.

SOUNDS LIKE “Is taste the new moat?”

INSTEAD OF “Having good taste might be the thing that sets you apart from the competition”

12 across · GROUNDBREAKING

The flagship of the breathless hype family, using superlatives as a substitute for proof: revolutionary, cutting-edge, innovative, game-changing, unprecedented. This default promotional tone that treats an ordinary product, idea, or update as unusually important. These words signal excitement, but often fail to give proof. AI-generated text commonly leans on these kinds of buzzword to make prose sound energetic and authoritative.

SOUNDS LIKE “Our groundbreaking new feature is a game-changer for productivity.”

INSTEAD OF “You can export to a spreadsheet now, which people have asked for since 2021.”

14 across · SYNERGY

This corporate cliché is conveniently vague and can be used to imply action without naming the actual mechanism or result. It suggests that collaboration will create value without committing the writer to a measurable claim or an accountable outcome, becoming a polished substitute for an explanation. Just like the words ‘ecosystem’ and ‘holistic’, it’s a low-risk filler. Corporate jargon has long worked this way; AI is merely reproducing an established human habit.

SOUNDS LIKE “The merger will unlock powerful synergies across both organisations.”

INSTEAD OF “They’ll share one warehouse and close the smaller one.”

1 down · QUIETLY

Frequently noticed in AI prose, especially phrases like “quietly reshaping,” “quietly transforming,” or “quietly becoming” because it creates an intriguing sense of momentum without having to identify who acted, when, or what changed.

SOUNDS LIKE “The company has quietly become one of the biggest landlords in the country.”

INSTEAD OF “The company owns 40,000 homes and almost nobody has heard of it.”

NO. 3 — THE CHAT-BACK

These are the most over-used conversational responses that models use when giving their output. When chatting with an AI assistant, it’s common for models to start by acknowledging, validating, and amplifying your request with “Absolutely” or “You’re exactly right.” But there are several other tics, patterns, and preferences and once you see them, it’s hard to un-see them. Different models have their own personality as well, so some of these you’ll see more often than not, depending on which you’re using. To varying degrees though, each model rewards rapport over efficiency. So whenever you’re ready we can jump right in.

across

3 Dots, dashes, or symbols used to break information into a list (7)

5 Extra help volunteered after the answer, unasked (9)

7 Portable lunch made with two pieces of bread and a filling. Also an AI answer with padding at both ends (8)

9 A strong “yes”, used to affirm what someone has said (10)

10 Two-word compliment paid to your question (13)

12 A wall or structure that supports the weight above it, AI uses it to mean “crucial” (11)

down

1 Small pictographs sprinkled through ChatGPT replies (6)

2 A mix of low quality wet food waste fed to pigs, now a common name for cheap AI-made content produced in bulk (4)

4 Being overly flattering or fawning (10)

6 Naming and affirming a user’s emotion to appear empathetic (10)

8 Everyday reasoning that keeps a model from suggesting the absurd (11)

11 Being correct, as in YOU’RE _____. And a particular tic of Claude Code (5)

THE ANSWERS, AND WHY

Every entry, with the note the website shows once you solve it – what the habit is, what it sounds like, and what someone would have written instead.

NO.1 — THE WORDS

2 across · GENUINELY

Similarly ‘honestly’, ‘truly’, and ‘really’ are used to add warmth and emphasis. ‘This is good’ becomes ‘this is genuinely good’. The extra word is a conversational sincerity marker and can make AI prose sound candid, helpful, and personally invested.

SOUNDS LIKE “Honestly, this is a genuinely exciting development that I truly believe matters.”

INSTEAD OF “This is a big deal, and I think it will hold up.”

5 across · UTILISE (or UTILIZE)

Here is the clearest example of AI preferring the longer version of a word instead of a more natural or colloquial word. Over time the longer word won thousands of small contests and became the default.

SOUNDS LIKE “Please utilise the attached template to facilitate the process.”

INSTEAD OF “Use the form attached. It takes two minutes.”

8 across · TAPESTRY

AI often uses this ornate noun as a shortcut for “a complex mixture of connected things.” It is an attractive, high-register metaphor making an ordinary point sound rich, thoughtful, and literary without requiring the model to name the actual parts or relationships.

SOUNDS LIKE “The local music scene is part of the rich tapestry of British culture.”

INSTEAD OF “Local venues, radio stations, promoters, and musicians have shaped British music for decades.”

9 across · DELVE

The use of this word rose 1,375% between 2020 and 2024. A 2025 analysis found that this word (and others strongly associated with ChatGPT-style prose like underscore, meticulous, and pivotal) began declining after being publicly identified as AI tells. This is partially down to human editing, like the ‘em dash’, once writers become aware of an AI trope they seek them out and remove or replace these tells. Newer models are also less likely to use some older giveaways.

SOUNDS LIKE “In this article we delve into the intricacies of remote work.”

INSTEAD OF “This article looks at remote work.”